

Module 2: Supervised Learning

Jianhui Gao, Siyue Yang, and Jessica Gronsbell

31/05/2022

```
# If a package is installed, it will be loaded. If any
## are not, the missing package(s) will be installed
## from CRAN and then loaded.

## First specify the packages of interest
packages <- c(
  "dplyr", "PheCAP", "glmnet", "randomForestSRC", "PheNorm",
  "MAP", "pROC", "mltools", "data.table", "ggplot2", "parallel"
)

## Now load or install&load all
package.check <- lapply(
  packages,
  FUN = function(x) {
    if (!require(x, character.only = TRUE)) {
      install.packages(x, dependencies = TRUE)
      library(x, character.only = TRUE)
    }
  }
)

# load environment from example 1
load("environment.RData")
```

Prepare data for algorithm development

- Split data into training and testing set
- Training 106(60%), Testing 75(40%)

```
data <- PhecapData(PheCAP::ehr_data, "healthcare_utilization", "label", 75,
  patient_id = "patient_id", seed = 123
)

# Data with non-missing labels
labeled_data <- ehr_data %>% dplyr::filter(!is.na(label))

# All Features
all_x <- ehr_data %>% dplyr::select(
  starts_with("COD"), starts_with("NLP"),
  starts_with("main"), healthcare_utilization
)
```

```

health_count <- ehr_data$healthcare_utilization

# Training Set
train_data <- ehr_data %>% dplyr::filter(patient_id %in% data$training_set)
train_x <- train_data %>%
  dplyr::select(
    starts_with("COD"), starts_with("NLP"),
    starts_with("main"), healthcare_utilization
  ) %>%
  as.matrix()
train_y <- train_data %>%
  dplyr::select(label) %>%
  pull()

# Testing Set
test_data <- ehr_data %>% dplyr::filter(patient_id %in% data$validation_set)
test_x <- test_data %>%
  dplyr::select(
    starts_with("COD"), starts_with("NLP"),
    starts_with("main"), healthcare_utilization
  ) %>%
  as.matrix()
test_y <- test_data %>%
  dplyr::select(label) %>%
  pull()

```

Penalized logistic regression

- Fit LASSO and Adaptive LASSO(ALASSO)

```

# Choose best lambda using CV
beta.lasso <- lasso_fit(
  x = train_x, y = train_y,
  tuning = "cv", family = "binomial"
)

# Features Selected
names(beta.lasso[abs(beta.lasso) > 0])[-1]

## [1] "NLP93" "NLP104" "NLP304"
## [4] "main_NLP" "main_ICDNLP" "healthcare_utilization"

# prediction on testing set
y_hat.lasso <- linear_model_predict(
  beta = beta.lasso, x = test_x,
  probability = TRUE
)

# Fit Adaptive LASSO
beta.alasso <- adaptive_lasso_fit(
  x = train_x, y = train_y,
  tuning = "cv", family = "binomial"
)
y_hat.alasso <- linear_model_predict(
  beta = beta.alasso, x = test_x,

```

```

    probability = TRUE
  )

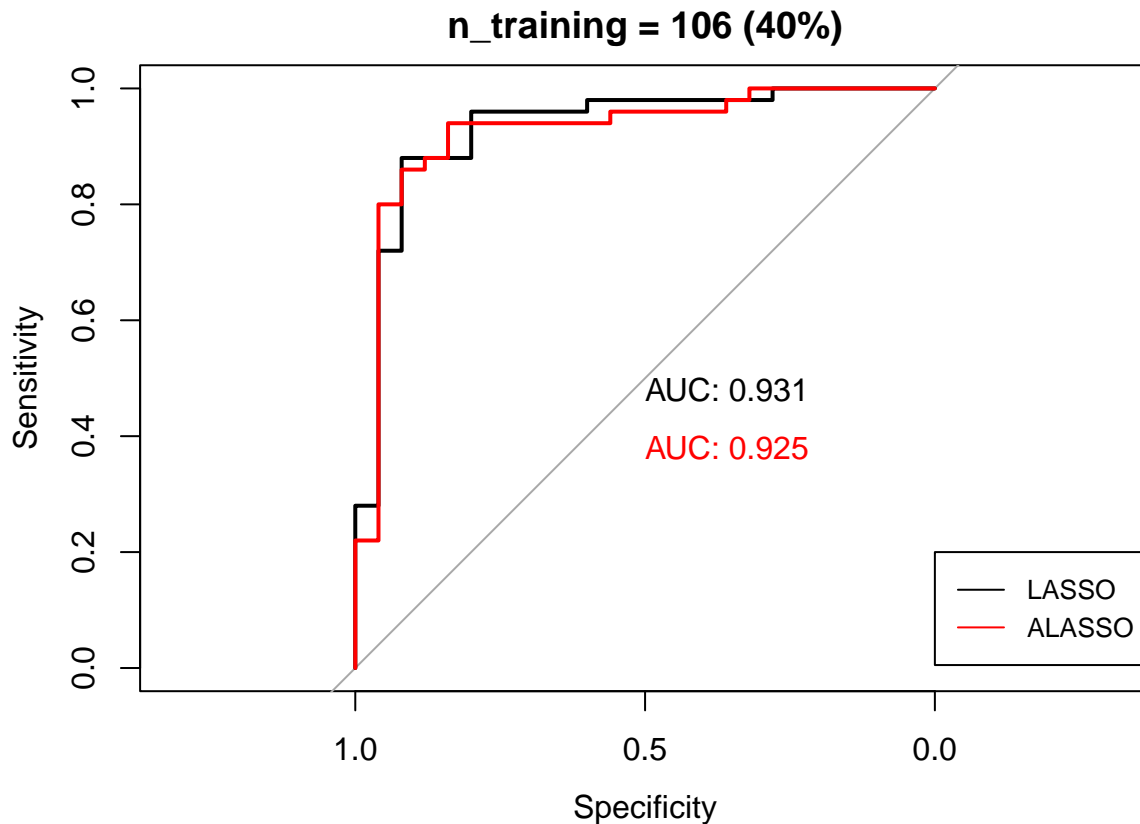
  # Features Selected
  names(beta.lasso[abs(beta.lasso) > 0])[-1]

## [1] "NLP304"                "main_NLP"                "healthcare_utilization"

roc.lasso <- roc(test_y, y_hat.lasso)
roc.lasso <- roc(test_y, y_hat.lasso)

plot(roc.lasso,
     print.auc = TRUE, main = "n_training = 106 (40%)")
)
plot(roc.lasso,
     print.auc = TRUE, col = "red", add = TRUE, print.auc.y = 0.4
)
legend(0, 0.2,
     legend = c("LASSO", "ALASSO"), col = c("black", "red"),
     lty = 1, cex = 0.8
)

```



```

roc_full.lasso <- get_roc(y_true = test_y, y_score = y_hat.lasso)
head(roc_full.lasso, 10)

```

```

##      cutoff  pos.rate  FPR    TPR    PPV    NPV    F1
## [1,] 0.9468465 0.006666667 0.00 0.1306667 1.0000000 0.3651412 0.2311321
## [2,] 0.9140868 0.066666667 0.00 0.2053333 1.0000000 0.3861998 0.3407080

```

```
## [3,] 0.8813270 0.193333333 0.02 0.2800000 0.9655172 0.4049587 0.4341085
## [4,] 0.8781599 0.200000000 0.04 0.3900000 0.9512195 0.4403670 0.5531915
## [5,] 0.8749928 0.206666667 0.04 0.5000000 0.9615385 0.4897959 0.6578947
## [6,] 0.8130579 0.413333333 0.04 0.6100000 0.9682540 0.5517241 0.7484663
## [7,] 0.7511229 0.500000000 0.06 0.7200000 0.9600000 0.6266667 0.8228571
## [8,] 0.7445305 0.506666667 0.08 0.7600000 0.9500000 0.6571429 0.8444444
## [9,] 0.7379380 0.513333333 0.08 0.8000000 0.9523810 0.6969697 0.8695652
## [10,] 0.6905611 0.586666667 0.08 0.8400000 0.9545455 0.7419355 0.8936170

roc_full.lasso <- get_roc(y_true = test_y, y_score = y_hat.lasso)
head(roc_full.lasso, 10)
```

```
##          cutoff    pos.rate  FPR      TPR      PPV      NPV      F1
## [1,] 0.9468465 0.006666667 0.00 0.1306667 1.0000000 0.3651412 0.2311321
## [2,] 0.9140868 0.066666667 0.00 0.2053333 1.0000000 0.3861998 0.3407080
## [3,] 0.8813270 0.193333333 0.02 0.2800000 0.9655172 0.4049587 0.4341085
## [4,] 0.8781599 0.200000000 0.04 0.3900000 0.9512195 0.4403670 0.5531915
## [5,] 0.8749928 0.206666667 0.04 0.5000000 0.9615385 0.4897959 0.6578947
## [6,] 0.8130579 0.413333333 0.04 0.6100000 0.9682540 0.5517241 0.7484663
## [7,] 0.7511229 0.500000000 0.06 0.7200000 0.9600000 0.6266667 0.8228571
## [8,] 0.7445305 0.506666667 0.08 0.7600000 0.9500000 0.6571429 0.8444444
## [9,] 0.7379380 0.513333333 0.08 0.8000000 0.9523810 0.6969697 0.8695652
## [10,] 0.6905611 0.586666667 0.08 0.8400000 0.9545455 0.7419355 0.8936170
```

Different train size

- randomly sample training size = 50, 70, 90
- rest as testing set
- repeat 600 times

```
start <- Sys.time()
auc_supervised <- validate_supervised(
  dat = labeled_data, nsim = 600,
  n.train = c(50, 70, 90)
)
end <- Sys.time()

end - start
```

```
## Time difference of 3.267942 mins
```

```
# median AUC
apply(auc_supervised, 2, median)
```

```
## n=50,LASSO n=70,LASSO n=90,LASSO n=50,ALASSO n=70,ALASSO n=90,ALASSO
## 0.8670982 0.8789683 0.8907670 0.8673935 0.8736602 0.8855655
```

```
# SE
apply(auc_supervised, 2, sd)
```


```
## n=50,LASSO n=70,LASSO n=90,LASSO n=50,ALASSO n=70,ALASSO n=90,ALASSO
## 0.07197811 0.05588511 0.05184181 0.07300341 0.05871336 0.05415953
```

```
par(mfrow = c(1, 3))
boxplot(auc_supervised %>% select(starts_with("n=50")),
  ylim = c(0.5, 1),
  names = c("LASSO", "ALASSO"), main = "n=50"
```

```

)
boxplot(auc_supervised %>% select(starts_with("n=70")),
  ylim = c(0.5, 1),
  names = c("LASSO", "ALASSO"), main = "n=70"
)
boxplot(auc_supervised %>% select(starts_with("n=90")),
  ylim = c(0.5, 1),
  names = c("LASSO", "ALASSO"), main = "n=90"
)

```

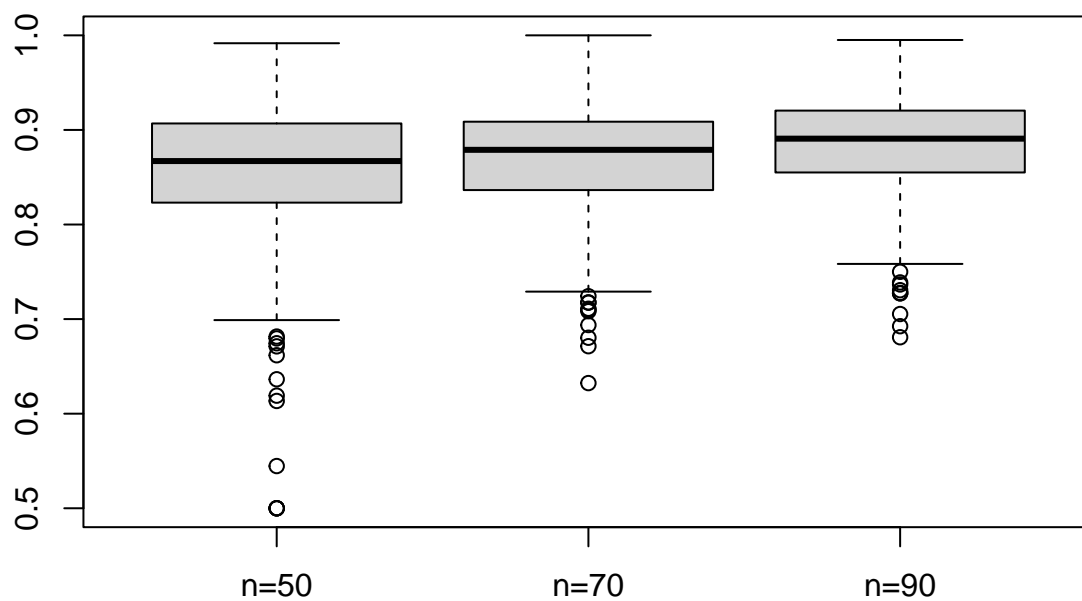


```

boxplot(auc_supervised[, 1:3],
  ylim = c(0.5, 1),
  names = c("n=50", "n=70", "n=90"), main = "LASSO"
)

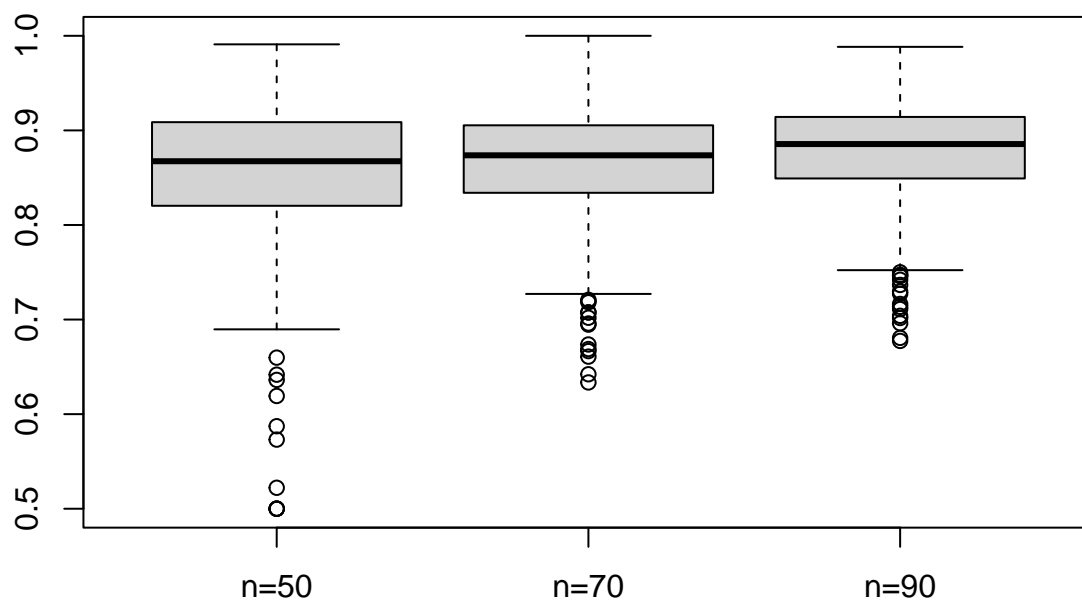
```

LASSO



```
boxplot(auc_supervised[, 4:6],  
        ylim = c(0.5, 1),  
        names = c("n=50", "n=70", "n=90"), main = "LASSO"  
)
```

ALASSO



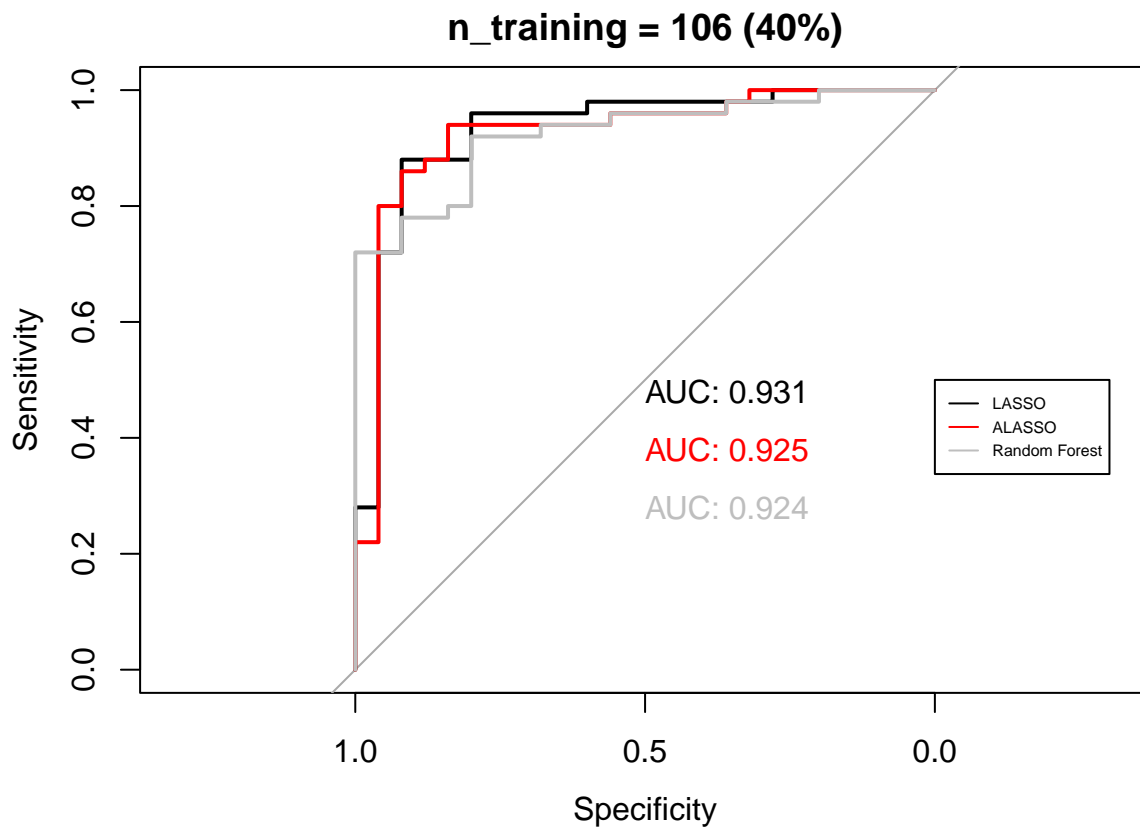
Appendix

Random Forest

```
model_rf <- rfsrc(y ~ ., data = data.frame(y = train_y, x = train_x))
y_hat.rf <- predict(model_rf, newdata = data.frame(x = test_x))$predicted

roc.rf <- roc(test_y, y_hat.rf)

plot(roc.lasso,
     print.auc = TRUE, main = "n_training = 106 (40%)")
)
plot(roc.lasso,
     print.auc = TRUE, col = "red", add = TRUE, print.auc.y = 0.4
)
plot(roc.rf,
     print.auc = TRUE, col = "grey", add = TRUE, print.auc.y = 0.3
)
legend(0, 0.5,
     legend = c("LASSO", "ALASSO", "Random Forest"), col = c("black", "red", "grey"),
     lty = 1, cex = 0.5
)
```



SVM

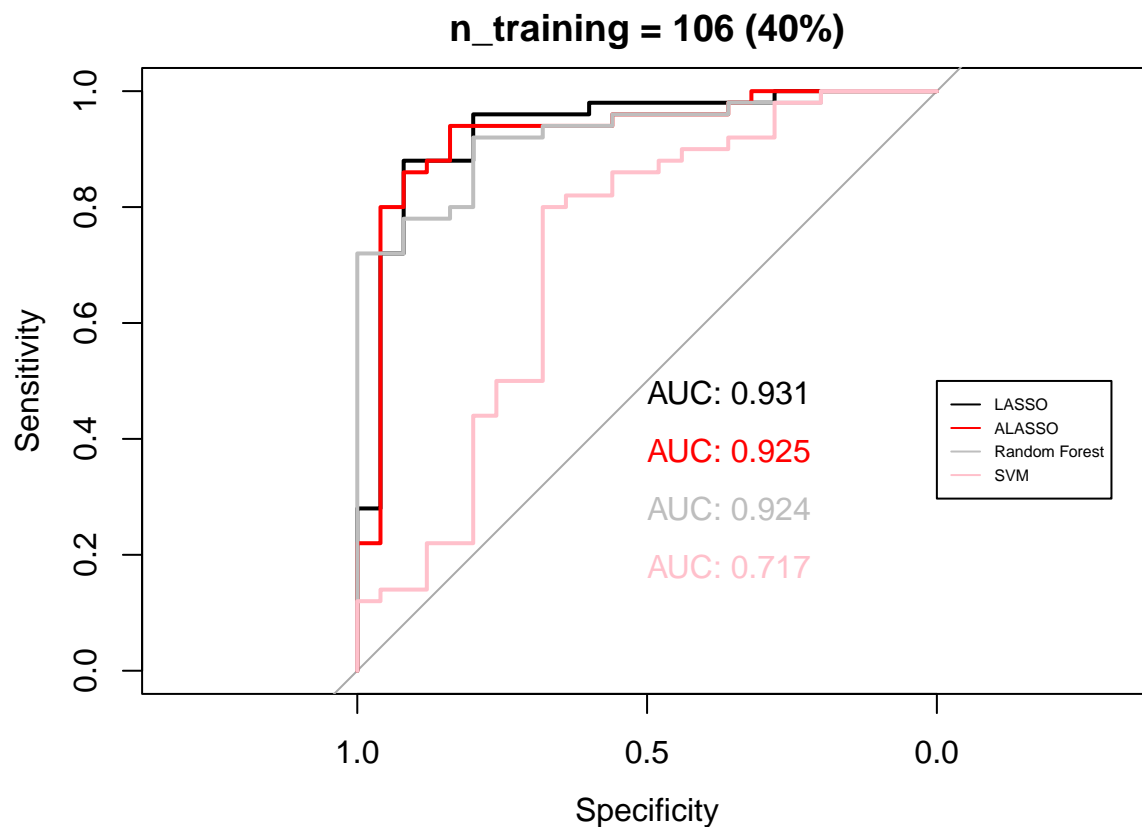
```
model_svm <- SVMmaj::svmmaj(X = train_x, y = train_y)
y_hat.svm <- predict(model_svm, test_x)

roc.svm <- roc(test_y, y_hat.svm)
```

```

plot(roc.lasso,
     print.auc = TRUE, main = "n_training = 106 (40%)")
)
plot(roc.lasso,
     print.auc = TRUE, col = "red", add = TRUE, print.auc.y = 0.4
)
plot(roc.rf,
     print.auc = TRUE, col = "grey", add = TRUE, print.auc.y = 0.3
)
plot(roc.svm,
     print.auc = TRUE, col = "pink", add = TRUE, print.auc.y = 0.2
)
legend(0, 0.5,
      legend = c("LASSO", "ALASSO", "Random Forest", "SVM"),
      col = c("black", "red", "grey", "pink"),
      lty = 1, cex = 0.5
)

```



Validation

```

start <- Sys.time()
auc_rfandsvm <- validate_svmandrf(dat = labeled_data, nsim = 600)
end <- Sys.time()
end - start
saveRDS(auc_rfandsvm, "appendix.rds")

```



```
auc_rfandsvm <- readRDS("appendix.rds")
```

```
# median AUC
```

```
apply(auc_rfandsvm, 2, median)
```

```
##      n=50,rf      n=70,rf      n=90,rf      n=50,svm      n=70,svm      n=90,svm
```

```
## 0.8701826 0.8873775 0.9008419 0.7179279 0.7540064 0.7848541
```

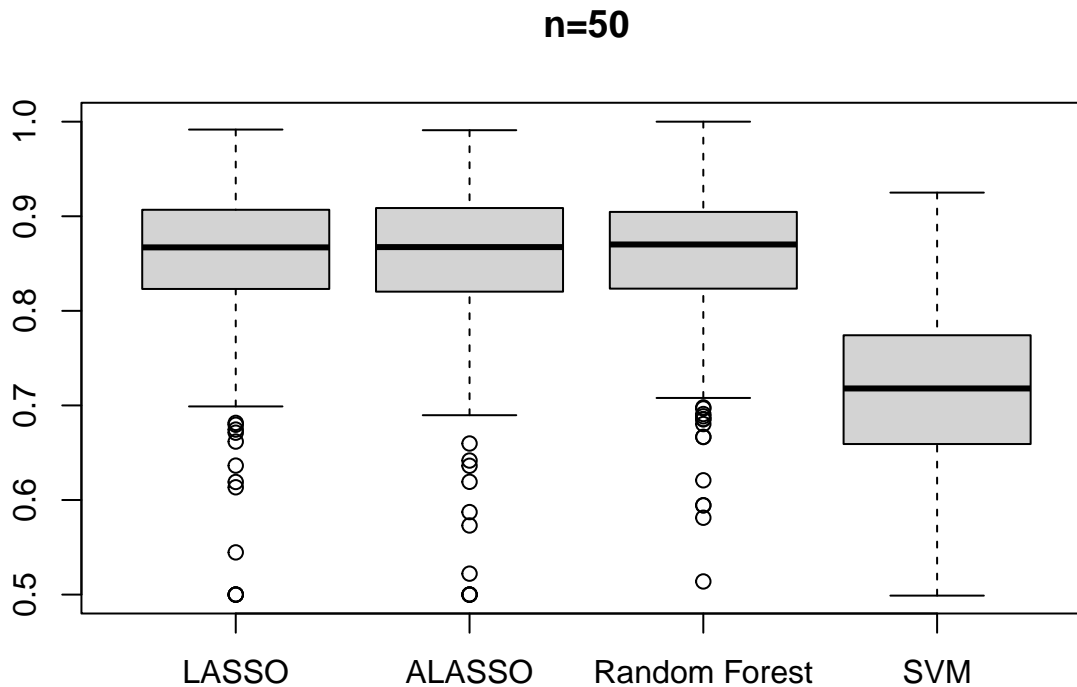
```
# SE
```

```
apply(auc_rfandsvm, 2, sd)
```

```
##      n=50,rf      n=70,rf      n=90,rf      n=50,svm      n=70,svm      n=90,svm
```


```
## 0.06771061 0.05256562 0.04799935 0.08221553 0.08092279 0.07128458
```

```
boxplot(cbind(auc_supervised, auc_rfandsvm) %>% select(starts_with("n=50")),
        ylim = c(0.5, 1),
        names = c("LASSO", "ALASSO", "Random Forest", "SVM"), main = "n=50"
)
```



```
boxplot(cbind(auc_supervised, auc_rfandsvm) %>% select(starts_with("n=70")),
        ylim = c(0.5, 1),
        names = c("LASSO", "ALASSO", "Random Forest", "SVM"), main = "n=70"
)
```

n=70



```
boxplot(cbind(auc_supervised, auc_rfandsvm) %>% select(starts_with("n=90")),  
        ylim = c(0.5, 1),  
        names = c("LASSO", "ALASSO", "Random Forest", "SVM"), main = "n=90"  
)
```

n=90

