

Module 2: Supervised learning

Split data into train and test

```
dim(train_x)
```

```
## [1] 106 587
```

```
length(train_y)
```

```
## [1] 106
```

```
dim(test_x)
```

```
## [1] 75 587
```

```
length(test_y)
```

```
## [1] 75
```

LASSO logistic regression

```
# Choose best lambda using CV.
```

```
beta_lasso <- lasso_fit(  
  x = log(train_x + 1),  
  y = train_y,  
  tuning = "cv",  
  family = "binomial"  
)
```

```
# Features Selected.
```

```
names(beta_lasso[abs(beta_lasso) > 0])[-1]
```

```
## [1] "COD2"      "COD10"     "NLP1"      "NLP17"     "NLP56"     "NLP82"  
## [7] "NLP93"     "NLP104"    "NLP118"    "NLP130"    "NLP144"    "NLP164"  
## [13] "NLP172"    "NLP193"    "NLP199"    "NLP222"    "NLP231"    "NLP265"  
## [19] "NLP274"    "NLP280"    "NLP297"    "NLP299"    "NLP346"    "NLP362"  
## [25] "NLP375"    "NLP382"    "NLP396"    "NLP401"    "NLP409"    "NLP435"  
## [31] "NLP451"    "NLP462"    "NLP488"    "NLP533"    "NLP536"    "NLP552"  
## [37] "NLP568"    "main_NLP"
```

ALASSO logistic regression

```
# Fit Adaptive LASSO.
beta_lasso <- adaptive_lasso_fit(
  x = log(train_x + 1),
  y = train_y,
  tuning = "cv",
  family = "binomial"
)
```

```
# ALASSO features selected.
beta_lasso[!beta_lasso == 0][-1]
```

```
##      NLP56      NLP93      NLP104      NLP118      NLP222      NLP231      NLP265
##  0.1966447 -1.0538342 -1.7011315 -1.5489010 -2.0758094  0.3598780 -0.9584738
##      NLP280      NLP297      NLP299      NLP409      NLP536      main_NLP
##  0.6256635 -0.2093127  1.0106695  0.4019735  0.1038460  1.4248803
```

```
# LASSO features selected.
beta_lasso[!beta_lasso == 0][-1]
```

```
##      COD2      COD10      NLP1      NLP17      NLP56      NLP82
## -0.07891435 -0.07964064 -0.15656996 -0.10698323  0.43476973 -0.14774013
##      NLP93      NLP104      NLP118      NLP130      NLP144      NLP164
## -0.95721897 -1.14198338 -0.83985826 -0.02971022 -0.39607669 -0.13824534
##      NLP172      NLP193      NLP199      NLP222      NLP231      NLP265
##  0.11876041  0.11493486 -0.16297872 -2.01541309  0.40654328 -0.84088955
##      NLP274      NLP280      NLP297      NLP299      NLP346      NLP362
## -0.17839805  0.62463549 -0.54371389  0.86087307 -0.40862069  0.17883546
##      NLP375      NLP382      NLP396      NLP401      NLP409      NLP435
##  0.79214450 -0.47973944 -0.08726960 -0.17450935  0.53175298  0.20241840
##      NLP451      NLP462      NLP488      NLP533      NLP536      NLP552
##  0.61949264 -0.24987822  0.46166193 -0.37801422  0.53979607  0.04623370
##      NLP568      main_NLP
##  0.40970337  1.28008994
```

Get model predictions + ROC curve


```
# Prediction on testing set (LASSO).  
y_hat_lasso <- linear_model_predict(  
  beta = beta_lasso,  
  x = log(test_x + 1),  
  probability = TRUE  
)
```

```
# Prediction on testing set (ALASSO).  
y_hat_lasso <- linear_model_predict(  
  beta = beta_lasso,  
  x = log(test_x + 1),  
  probability = TRUE  
)
```

```
roc_lasso <- roc(test_y, y_hat_lasso)  
roc_lasso <- roc(test_y, y_hat_lasso)
```

LASSO vs. ALASSO

The operating receiver characteristic (ROC) curve



LASSO vs. ALASSO at $FPR = 0.10$

```
roc_full_lasso <- get_roc(y_true = test_y, y_score = y_hat_lasso) %>% data.frame()
get_roc_parameter(0.1, roc_full_lasso)
```

```
##      cutoff pos.rate FPR  TPR      PPV      NPV      F1
## 1 0.8573308 0.5933333 0.1 0.84 0.9438202 0.7377049 0.8888889
```

```
roc_full_lasso <- get_roc(y_true = test_y, y_score = y_hat_lasso) %>% data.frame()
get_roc_parameter(0.1, roc_full_lasso)
```

```
##      cutoff pos.rate FPR  TPR      PPV      NPV      F1
## 1 0.9228109 0.5133333 0.1 0.72 0.9350649 0.6164384 0.8135593
```


LASSO vs. ALASSO with different training set size

- ▶ Randomly sample training size = 50, 70, 90
- ▶ Use the remaining data as the test set
- ▶ Repeat 600 times

```
auc_supervised <- validate_supervised(  
  dat = labeled_data,  
  nsim = 600,  
  ntrain = c(50, 70, 90)  
)
```


LASSO vs. ALASSO with different training set size

Area under the ROC curve (AUC) from 600 simulations



Random Forest and SVM

```
# Random forest.
```

```
model_rf <- rfsrc(y ~ ., data = data.frame(y = train_y, x = train_x))  
y_hat_rf <- predict(model_rf, newdata = data.frame(x = test_x))$predicted  
roc_rf <- roc(test_y, y_hat_rf)
```

```
# SVM.
```

```
model_svm <- SVMmaj::svmmaj(X = train_x, y = train_y)  
y_hat_svm <- predict(model_svm, test_x)  
roc_svm <- roc(test_y, y_hat_svm)
```

ROC curves

The operating receiver characteristic (ROC) curve

